

A New Approach to Analyzing Robin Hood Hashing

Michael Mitzenmacher*

Abstract

Robin Hood hashing is a variation on open addressing hashing designed to reduce the maximum search time as well as the variance in the search time for elements in the hash table. While the case of insertions only using Robin Hood hashing is well understood, the behavior with deletions has remained open. Here we show that Robin Hood hashing with random hash functions can be analyzed under the framework of finite-level finite-dimensional jump Markov chains. This framework allows us to re-derive some past results for the insertion-only case with some new insight, as well as provide new analyses for a standard deletion model, where we alternate between deleting a random old key and inserting a new one.

1 Introduction

Robin Hood hashing is a variation on open addressing hashing designed to reduce the maximum search time as well as the variance in the search time for elements in the hash table. Here we are interested in the setting where the probe sequences are random. We briefly describe the setup, starting with a setting with insertions only. We have a hash table with n cells, and $m = \lceil \alpha n \rceil$ keys to place in the table. We refer to α as the load of the table; generally, we assume αn is an integer henceforth. Each key K_i has an associated infinite probe sequence K_{ij} , with $j \geq 1$, where the K_{ij} are independently and uniformly distributed over $[0, n - 1]$. Equivalently, the K_{ij} are determined by a random hash function h , where for a keyspace K the hash function has the form $h : K \times \mathbb{N} \rightarrow [0, n - 1]$.¹ Each key will be placed according to a position in its probe sequence. If the i th element is placed in cell K_{ij} , and there is no $j' < j$ such that $K_{ij'} = K_{ij}$, we shall say that the *age* of the key is j . If we use the standard search process for a key, by which we mean sequentially examining cells according to the

probe sequence, the age of a key in the table corresponds to the number of cells that must be searched to find it. We assume that we keep track of the age of the oldest key in the table. In the standard search process, one determines that a key not in the table is not present by sequentially examining cells according to the probe sequence until either an empty cell is found, or one has found that the key being searched for must be older than oldest key in the table. An empty cell provides a witness that the key is not in the table. We refer to a search for a key not in the system as an unsuccessful search.

For the insertion of keys in the table, we may think of the keys as being *placed* sequentially, using the probe sequence in the following manner. If K_{i1} is empty when the i th key is inserted into the table, the key is readily placed at cell K_{i1} . Otherwise, there is a collision, and a collision resolution strategy is required. The main point of Robin Hood hashing is that it *resolves collisions in favor of the key with the larger age*; the key with the smaller age must continue sequentially through its probe sequence. Notice that, under Robin Hood hashing, a placed key will be displaced by the key currently being placed if the placed key's age is smaller. In this case the placed key is moved from its current cell and becomes the item to be placed, consequentially increasing its age. Other standard conflict resolution mechanisms are first come first served and last come first served. By favoring more aged keys, Robin Hood hashing aims to reduce the maximum search time required.

Most of the results for Robin Hood hashing appear in the thesis of Celis [2], who provides a number of theoretical and empirical results. (See also [3].) The following results are especially worth mentioning. First, when there are only insertions, Celis analyzes the asymptotic behavior of Robin Hood hashing (in the infinite limit setting) for loads $\alpha < 1$. We describe this result further in Section 3.2. Second, Celis shows that the total expected insertion cost in terms of the number of probes evaluated by the standard insertion process – or equivalently the average age of keys in the table – is the same for a class of “oblivious” collision resolution strategies that do not make use of knowledge about the future values in the probe sequences and that include Robin Hood hashing (as well as first come first served

*School of Engineering and Applied Sciences, Harvard University. This work was supported by NSF grants CCF-0915922, IIS-0964473, CNS-1228598, and CCF-1320231. Part of this work was done while visiting Microsoft Research New England.

¹Alternatively, we could have each probe sequence be a random permutation of $[0, n - 1]$ for each key; for our purposes, the two models are essentially equivalent, and we use the random hash function model as it is easier to work with.

and last come first served). Third, Devroye, Morin, and Viola have shown that for $\alpha < 1$ the maximum search time for Robin Hood hashing is (upper and lower) bounded by $\log_2 \log_2 n \pm O(1)$ with probability $1 - o(1)$, where the $O(1)$ terms depend on α [4]. This double-logarithmic behavior also occurs with quite different hashing schemes based on the power of multiple choices [1, 12]. Finally, we note that Robin Hood hashing has been also studied extensively in the setting of linear probing schemes [5, 15, 16].

Here we provide a new approach for analyzing Robin Hood hashing, based on a fluid limit analysis utilizing differential equations. An interesting aspect of our analysis is that it uses an additional level parameter, corresponding to a faster-moving Markov process (tracking the age of current key being placed) beyond the larger-scale Markov process (tracking the distribution of ages in the table). This type of analysis was previously used to study load balancing schemes with memory [7, 11]. Our analysis allows us to re-derive previous results for Robin Hood hashing, such as the asymptotic behavior for loads $\alpha < 1$, while also providing some additional novelty, such as concentration bounds for finite n . We also re-derive the high probability upper bound on the maximum search time for Robin Hood hashing of $\log_2 \log_2 n + O(1)$ of [4], with what we suggest is a simpler and more intuitive proof. More importantly, our approach is amenable to studying Robin Hood hashing with deletions of random keys, an area that lacked a theoretical framework for analysis previously. We study the deletion scheme proposed by Celis in [2] under the setting of random deletions of keys and new keys being inserted (maintaining a constant load α), and suggest and analyze an alternative deletion scheme that is simpler for practical implementations.

Before beginning, we remark that practical use is not our main motivation for studying the Robin Hood hashing variant we examine here (although we have seen some suggestions that it is still used on occasion). Robin Hood hashing can require substantially more randomness than many other hashing schemes (such as cuckoo hashing or other multiple-choice hashing schemes), and practical considerations such as cache performance and prefetching suggest that one would prefer to use the linear probing variant in almost all settings. Our motivation instead is in understanding this classical and combinatorially simple-seeming hashing scheme, as well as in the techniques that can be used to analyze it. In particular, the double logarithmic bound on the search time requires some non-trivial additional technical work beyond the standard layered induction approach, and our analysis of performance with deletions appears entirely new.

In what follows, we provide background on the fluid limit approach we use here. We then study Robin Hood hashing in the setting of insertions only under this framework, and subsequently move on to examining how to analyze settings with random deletions. We note that our work includes extensive simulations that demonstrate the accuracy of our approach. We put these results in appendices.

2 Limiting Framework

For our limiting framework, we can work in the setting of finite-level finite-dimensional jump Markov chains. Here we roughly follow the exposition of [11]; further development can be found in [14]. Our discussion here is brief, and may be skipped by the uninterested reader willing to accept the more intuitive explanations that follow. We suspect this methodology should be useful for studying other hashing variations.

In our setting, a chain with D dimensions and L levels will have the state space $\mathbb{R}^D \times \{1, 2, \dots, L\}$. The state can be represented as a $D+L$ -tuple in the natural way as follows: $(\bar{x}; m) = (x_1, \dots, x_D; 0, \dots, 1, \dots, 0)$, where a 1 in position $D+m$, $1 \leq m \leq L$, represents that the system is in level m . When in state $(\bar{x}; m)$ the system can make $\zeta(m)$ possible different jumps. Here we describe only unit jumps based on unit vectors, which suffices for our main application, but more general jumps are possible. The process jumps to state $(\bar{x} + \bar{e}_i(m); g_i(m)) = (\bar{x} + \bar{e}_i(m); 0, \dots, 1, \dots, 0)$ with rate $\nu_i(\bar{x}; m)$, for $1 \leq i \leq \zeta(m)$, where the 1 is in position $D + g_i(m)$. Here $\bar{e}_i(m)$ is a unit vector in one of the D dimensions, and $g_i(m)$ is the (new) level associated with the i th of the $\zeta(m)$ possible jumps; note that $g_i(m)$ might itself be m , so that the level may not change. The high-level idea is that here we have an underlying finite-dimensional jump Markov process, but we also have an additional associated “level” process that may drive the transition rates of the primary jump Markov process.

The generator A of this Markov process, which operates on real valued functions $f : \mathbb{R}^{D+L} \rightarrow \mathbb{R}$, is defined as:

$$Af(\bar{x}; m) = \sum_{i=1}^{\zeta(m)} \nu_i(\bar{x}; m) [f(\bar{x} + \bar{e}_i(m); g_i(m)) - f(\bar{x}; m)].$$

We now consider a scaled version of this process, with scaling parameter n , where the rate of each transition is scaled up by a factor of n and the jump magnitude is scaled down by a factor of n . The state of this scaled system will be represented by $(\bar{s}_n; m) = (s_1, \dots, s_D; 0, \dots, \frac{1}{n}, \dots, 0)$ (with now a $1/n$ term in the position for level m). The associated jump vectors will be $(\frac{\bar{e}_i}{n}; 0, \dots, -\frac{1}{n}, \dots, \frac{1}{n}, \dots, 0)$, with corresponding rates $n\nu_i(\bar{s}_n; m)$ for $1 \leq i \leq \zeta(m)$. (Note that in the case where the level does not change, the $-\frac{1}{n}$ and $\frac{1}{n}$

jumps in the level should be interpreted as being in the same coordinate, so no change occurs.) The generator for the scaled Markov process is:

$$A_n f(\bar{s}_n; m) = \sum_{i=1}^{\zeta(m)} n \nu_i(\bar{s}_n; m) \{f[\bar{s}_n + \frac{\bar{e}_i(m)}{n}; g_i(m)] - f(\bar{s}_n; m)\}$$

The following theorem (Theorem 8.15 from [14]) describes the evolution of the typical path of the scaled Markov process in the limit as n grows large. The idea behind the theorem is that because the finite-level Markov chain reaches equilibrium in some finite time, for large enough n the approximation that the finite-level Markov chain is in equilibrium is sufficient to obtain Chernoff-like bounds.

THEOREM 2.1. *Under Conditions 2.1 and 2.2 below, for any given T and constant $\epsilon > 0$, there exist positive constants $C_1, C_2(\epsilon)$ and n_0 such that for all initial positions $\bar{s}^0 \in \mathbb{R}^D$, any initial level $m \in \{0, 1, \dots, L-1\}$, and any $n \geq n_0$,*

$$\Pr_{\bar{s}^0, m} \left(\sup_{0 \leq t \leq T} |\bar{s}_n(t) - \bar{s}_\infty(t)| > \epsilon \right) \leq C_1 \exp(-n C_2(\epsilon)),$$

where $\bar{s}_\infty(t)$ satisfies the following:

$$\begin{aligned} \frac{d}{dt} \bar{s}_\infty(t) &= \sum_{l=0}^L \Pr(m(t) = l) \sum_{i=1}^{\zeta(l)} \nu_i(\bar{s}_\infty; l) \bar{e}_i(l); \\ \bar{s}_\infty(0) &= \bar{s}^0. \end{aligned}$$

Here $\Pr(m(t) = l)$ is the equilibrium probability of the level-process being in level l given the state $\bar{s}_\infty(t)$.

CONDITION 2.1. *For any fixed value of $\bar{x} \in \mathbb{R}^D$, the Markov process evolving over the levels $\{1, \dots, L\}$ with transition rate $\nu_i(\bar{x}; m)$ of going to level $g_i(m)$ from level m , is ergodic.*

CONDITION 2.2. *The functions $\log \nu_i(\bar{x}; y)$ are bounded and Lipschitz continuous in \bar{x} for every y (where continuity is in all the D coordinates).*

A limitation of this approach is that it directly provides bounds only on the finite-dimensional version of the process. It is thus not immediate that one can obtain rigorous bounds on the maximum search time for Robin Hood of the form $\log_2 \log_2 n \pm O(1)$ as in [4] directly using this approach, as tracking $\log_2 \log_2 n$ dimensions takes us outside the finite-dimensional realm. Instead, one may use these results as useful intuition to guide non-limiting probabilistic arguments such as that derived in [4], as we show here. In return for this limitation, however, this approach provides simple and general means for generating rich, accurate numerical results that can aid in design and performance testing for real-world implementations.

3 Robin Hood Hashing with Insertions Only

3.1 Applying the Limiting Framework

We first describe the standard process, which corresponds to the unscaled process described above; we generally use the term unscaled process where the meaning is clear. Each time step corresponds to an attempted placement of a key, which can either be a new key, or a key that was not successfully placed at the last time step, or a key that was displaced by another key at the last time step.

To keep track of the state, we take advantage of the fact that keys are placed randomly into cells. Hence, for the state it suffices to track the number of cells holding keys for each age; their actual position does not matter. Each time step corresponds to an attempt to place a key. Note that the number of time steps here is *not* equal to the number of keys placed; placing a new key can take several time steps with Robin Hood hashing, and as discussed during that process the key being placed can take the place of another key which then has to be placed. Each such placement attempt represents a time step. As is often the case with hashing schemes, we find it more useful to look at the tails of the loads rather than the loads themselves. Therefore, we let $x_i(t)$ be the fraction of cells with a key with age at least i after t unscaled time steps. For our scaled version of the state, we let $s_i(t)$ be the fraction of cells holding a key of age at least i after tn key placements have been tried, so that $x_i(nt) = s_i(t)$. The level in our process will correspond to the age of key currently being placed. Fresh keys that are newly being inserted have age 1.

We note that, as described, the process is infinite-dimensional, in that we can consider the values s_i for all $i \geq 1$. Indeed, this is usually how we will think of the process, although as we show later we can “truncate” the system at any finite value of i , which can allow us to apply Theorem 2.1.

As a warm-up in understanding the scaled process, note that when the load of the table is z , so that zn cells contain a key, the number of time steps to place an element is geometrically distributed with mean $1/(1-z)$. Hence in the limiting scaled process \bar{s}_∞ , with the initial state being an empty table, the load will be α at time

$$\int_{z=0}^{\alpha} \frac{1}{1-z} dz = \ln \frac{1}{1-\alpha}.$$

That is, we run until the scaled time $\ln \frac{1}{1-\alpha}$, which corresponds to (in the unscaled process, asymptotically) $n \ln \frac{1}{1-\alpha}$ time steps. Alternatively, in the limiting scaled process, at time t , the load is $1 - e^{-t}$.

We now turn to understanding the level process, assuming that the state of the table is fixed according

to the values s_i . Again we find it useful to consider the tails. Thinking of the unscaled process, so t again refers to discrete time steps, let $p_i(t)$ be the probability that the age of the key being placed is at least i . Hence $p_1(t) = 1$ for all time steps. For $i > 1$, the key being placed at time $t + 1$ will have age at least i if and only if both the age of the key being placed at time t has age at least $i - 1$, and cell chosen by the probe sequence at time t has age at least $i - 1$. This is because, assuming an empty cell is not found, the younger of the keys will be the key being placed at the next time step. In equation form, we simply have

$$p_i(t + 1) = p_{i-1}(t)x_{i-1}(t).$$

In the scaled time setting, this can be written as

$$(3.1) \quad p_i(t + 1/n) = p_{i-1}(t)s_{i-1}(t).$$

Notice that the simplicity of this equation helps justify our decision to focus on variables that represent the tails of the loads.

Since s_1 is bounded by the final load α , at each step with probability at least $1 - \alpha$ a key is placed and the level returns to 1. Hence the Markov process over the levels is ergodic. Indeed, standard methods show that for any constant ϵ this Markov can be made ϵ -close in statistical distance to its equilibrium distribution after some corresponding constant number of placement steps. We re-emphasize the intuition; in the scaled process, the s_i values change significantly (that is, by $\Omega(1)$) only after $\Omega(1)$ scaled time steps (or $\Omega(n)$ unscaled time steps), while the Markov chain governing the p_i converges (arbitrarily closely) to its stationary distribution in $o(n)$ unscaled time steps. Hence, it makes sense in the limit to treat the p_i values as fixed in equilibrium given the s_i values.

It follows from Equation (3.1) that the equilibrium for the level process satisfies

$$(3.2) \quad p_i = p_{i-1}s_{i-1}$$

when we treat the s_i as fixed and we use p_i without the t to denote the equilibrium for the p_i values given the s_i values at that time.

With this we turn our attention to the limiting equations for the s_i in \bar{s}_∞ . Note that s_1 increases whenever a empty cell is hit. Hence

$$(3.3) \quad \frac{ds_1}{dt} = 1 - s_1.$$

Integrating, and using $s_1(0) = 0$, this gives $s_1(t) = 1 - e^{-t}$, matching our previous warm-up analysis. For s_i when $i > 1$, Equation (3.3) generalizes to

$$(3.4) \quad \frac{ds_i}{dt} = p_i(1 - s_i),$$

since a cell containing a key with age at least i is created whenever the age of the key being placed is at least i and the probe sequence finds either an empty cell or a cell containing a key with age less than i .

At any time t , let $\beta(t)$ be the corresponding load at that time. (Recall we use α for the “final” load.) We use β for $\beta(t)$ where the meaning is clear. Since $\beta(t) = 1 - e^{-t}$, we have

$$\frac{d\beta}{dt} = e^{-t} = 1 - \beta.$$

At the possible risk of confusion, but to avoid conversions back and forth, we use $\hat{s}_i(\beta)$ to represent s_i taken as a function of the load β instead of as a function of time. We have from the above that in the setting of the asymptotic limit $\hat{s}_i(\beta) = s_i(-\ln(1 - \beta))$. With the expressions for $\frac{d\beta}{dt}$ and $\frac{ds_i}{dt}$ we obtain the following form for $\hat{s}_i(\beta)$ as a function of β for $i \geq 1$.

$$(3.5) \quad \frac{d\hat{s}_i}{d\beta} = \frac{p_i(1 - \hat{s}_i)}{1 - \beta}.$$

Given our equations for p_i , we can substitute so that all equations are in terms of the s_i . Specifically, the equation for $\frac{ds_i}{dt}$ (or $\frac{d\hat{s}_i}{d\beta}$) depends only on values s_j with $j \leq i$. That is,

$$\begin{aligned} \frac{ds_i}{dt} &= p_i(1 - s_i) \\ &= p_{i-1}s_{i-1}(1 - s_i) \\ &= (1 - s_i) \prod_{j=1}^{i-1} s_j. \end{aligned}$$

The differential equations can therefore be solved numerically for s_i values up to any desired constant K . Moreover, we can truncate the infinite system of differential equations to a finite system by considering the equations for $\frac{ds_i}{dt}$ up to the constant K . Because of this, using the large deviation theory, we may formally state the following:

THEOREM 3.1. *For any fixed constant K and any constant $\alpha < 1$, for $i \leq K$, let $\hat{s}_i(\alpha)$ be the solution for the \hat{s}_i at final load α from the family of differential equations given by Equation (3.5) above. For $1 \leq i \leq K$, let $X_{i,n}$ be the random variable denoting the fraction of cells with keys of age at least i using Robin Hood hashing at final load α with n cells. Then for any $\epsilon > 0$, for sufficiently large n*

$$\Pr(|X_{i,n} - \hat{s}_i(\alpha)| > \epsilon) \leq C_1 \exp(-nC_2(\epsilon)),$$

where C_1 is a constant that depends on K and α , and $C_2(\epsilon)$ is a constant that depends on K , α , and ϵ .

Proof. The result follows from Theorem 2.1. While Theorem 2.1 is stated in terms of time instead of load, this difference is not consequential; the straightforward proof is given in an appendix. \square

3.2 Implications for the Age Distribution

In [2], Celis derives the age distribution under Robin Hood hashing by providing a recurrence. We demonstrate that this result also follows from our differential equations analysis. We note that we use a different notation; the following theorem corresponds to Theorem 3.1 of [2].

THEOREM 3.2. *In the asymptotic model for an infinite Robin Hood hash table with load factor β ($\beta < 1$), the fraction $\hat{s}_i(\beta)$ of cells containing keys of age at least i is given by*

$$(3.6) \quad \hat{s}_{i+1}(\beta) = 1 - (1 - \beta)e^{\sum_{j=1}^i \hat{s}_j(\beta)}.$$

Proof. As standard techniques can be used to show that our family of differential equations has a unique solution, we show that the recurrence of Equation (3.6) satisfies Equation (3.5). We first note the following useful fact:

$$\sum_{j=1}^i \frac{d\hat{s}_j}{d\beta} = \frac{1 - \prod_{j=1}^i \hat{s}_j}{1 - \beta}.$$

This follows easily from Equation (3.5) by induction, using that $p_i = \prod_{j=1}^{i-1} s_j = \prod_{j=1}^{i-1} \hat{s}_j$, as is easily derived from Equation (3.2).

Now taking the derivative of Equation (3.6) we find

$$\begin{aligned} \frac{d\hat{s}_{i+1}(\beta)}{d\beta} &= e^{\sum_{j=1}^i \hat{s}_j} - (1 - \beta) \left(\sum_{j=1}^i \frac{d\hat{s}_j}{d\beta} \right) e^{\sum_{j=1}^i \hat{s}_j} \\ &= \frac{1 - \hat{s}_{i+1}}{1 - \beta} - (1 - \hat{s}_{i+1}) \frac{1 - \prod_{j=1}^i \hat{s}_j}{1 - \beta} \\ &= \frac{\left(\prod_{j=1}^i \hat{s}_j \right) (1 - \hat{s}_{i+1})}{1 - \beta} \\ &= \frac{p_{i+1}(1 - \hat{s}_{i+1})}{1 - \beta}. \end{aligned}$$

Hence the recurrence of Equation (3.6) satisfies Equation (3.5) as claimed. \square

3.3 Implications for Maximum Age

We now show, following an approach established in [8], that the fact that the growth of maximum age grows double logarithmically in n appears as a natural consequence of the differential equations. As noted, the general large deviation results we apply only hold

for finite-dimensional systems, so their application can only apply up to constant ages. (Of course, one can choose a very large constant age, so our extension here is clearly primarily of theoretical interest.) Explicitly proving an $O(\log \log n)$ bound on the maximum can be accomplished by translating the differential equation argument to a layered induction argument, successively bounding the fraction of cells holding keys of age i for each i . While it does not appear motivated by the differential equations approach, a previous argument in roughly the layered induction style appears already in [4], formally providing a $\log \log n + O(1)$ bound. Their analysis is very different, however, as it does not make use of the underlying Markov chain directly, but bounds the behavior of what the authors call the “head” and the “belly” of the process over stages to resolve the collision that arises. Our goal in this analysis is to first show how the fluid limit analysis provides novel insight into how the doubly exponential decrease in the age distribution arises. We then use the fluid limit argument as a guide, leading to an alternative (and we believe somewhat simpler) layered induction proof for a $\log \log n + O(1)$ bound on the maximum age.

THEOREM 3.3. *In the asymptotic model for an infinite Robin Hood hash table with load factor $\alpha < 1$, for sufficiently large constants i , the fraction $\hat{s}_i(\alpha)$ of cells that contain keys of age at least i satisfies*

$$(3.7) \quad \hat{s}_i(\alpha) \leq c_1 c_2^{2^{i-c_3}}$$

for some constants $c_1, c_3 > 0$ and $c_2 < 1$ that may depend on α .

Proof. In what follows here let $u = -\ln(1 - \alpha)$. Let j be the smallest value such that $s_j(u) < u^{-1}$. Let $s_j(u) = \nu$ and $\nu \cdot u = \nu^* < 1$. We remark that while it may not be immediately clear that the s_j go to 0, it is shown in [4, Lemma 4] that the s_i have geometrically decreasing tails, so that j is in fact a constant. (Alternatively, since the p_j clearly have geometrically decreasing tails, it follows readily that the s_j do as well.) Now below we use the differential equations based on time, up until time u , so $t \leq u$. As $p_{j+1} \leq s_j$, and s_j is increasing over time,

$$\frac{ds_{j+1}(t)}{dt} = p_{j+1}(1 - s_{j+1}) \leq s_j \leq \nu.$$

To reach load α we run for time $u = \ln \frac{1}{1-\alpha}$, and hence

$$\hat{s}_{j+1}(\alpha) = s_{j+1}(u) \leq \nu \cdot u = \nu^*.$$

Inductively, we now find by the same argument that for $k \geq 1$,

$$p_{j+k} \leq (\nu^*)^{2^{k-1}-1} \nu \quad ; \quad s_{j+k} \leq (\nu^*)^{2^{k-1}},$$

and the result follows. That is,

$$p_{j+k} = p_j \prod_{\ell=0}^{k-1} s_{j+\ell} \leq \nu \prod_{\ell=1}^{k-1} s_{j+\ell} \leq (\nu^*)^{2^{k-1}-1} \nu,$$

where the last step follows from the inductive hypothesis. Further, as

$$\frac{ds_{j+k}(t)}{dt} \leq p_{j+k} \leq (\nu^*)^{2^{k-1}-1} \nu,$$

we have

$$\hat{s}_{j+k}(\alpha) \leq (\nu^*)^{2^{k-1}-1} \nu u = (\nu^*)^{2^{k-1}}.$$

The theorem follows. \square

As one might hope from previous work (e.g., [8]), this fluid limit argument can be transformed into a layered induction argument to prove a $\log \log n + O(1)$ upper bound on the maximum age, as we now show.

THEOREM 3.4. *Let M_n be the maximum age in a Robin Hood hash table with n cells and load factor $\alpha < 1$. There is a constant C depending only on α such that*

$$\lim_{n \rightarrow \infty} \Pr(M_n \geq \log \log n + C) = 0.$$

Proof. Instead of thinking about the Robin Hood hash table, we work with the Markov process we have been considering, where here we take the time t to be the number of unscaled time steps, $S_i(t)$ to be the number of cells holding a key of age at least i at time t (where we include the key being placed in the count if its age is at least i), $s_i(t)$ to be $S_i(t)/n$, and $P(t)$ to be the age of the key being placed at the t th time step. The idea of the proof is to start with a bound on $S_j(t)$ at the end of the process for some useful starting value j , and use this to bound the number of steps for which $P(t) \geq j+1$ over the course of the process. With this, we in turn bound the number of steps that yield a new key of age at least $j+1$ to bound $S_{j+1}(t)$ at the end of the process, providing our induction. This induction, which uses Chernoff bounds, breaks down at some age (where the Chernoff bound will no longer readily apply), at which point we explicitly have to carefully cap off the induction. We follow the intuition from our argument for the limiting system. Finally, we note for this proof we have not attempted to optimize the constant C .

Here we let $u = -\ln(1 - (\alpha + \epsilon_0))$ for a small constant $\epsilon_0 < (1 - \alpha)/2$. We run the process for up to nu steps; however, once αn keys are placed, we allow the process to stop, and no action is taken in the remaining time steps. (Recall that αn keys should be placed after $-\ln(1 - (\alpha))n + o(n)$ steps in expectation.) Let \mathcal{G} be

the event that αn keys have been placed after nu steps. Using standard martingale arguments, one can easily show $\Pr(\mathcal{G}) = 1 - o(1)$; we return to this point at the end of the analysis.

For convenience we take the case where $u > 1/2$, so $j \geq 2$; this suffices as smaller α and hence values of u will have strictly smaller load. After nu steps, the total age of all keys in the table is at most nu . Therefore the fraction of keys in the table with age at least $j = \lceil 16u^3 \rceil$ is at most $(4u)^{-2}$, so we deterministically have $s_j(nu) \leq (4u)^{-2}$. As a warm up, following the fluid limit argument, let X_t^{j+1} be 1 if $P(t) \geq j+1$ and 0 otherwise. For $P(t) \geq j+1$, the cell chosen at the previous step must have age at least j . For each such time step, the probability such a cell was chosen is dominated by an independent Bernoulli trial of probability $s_j(nu) \leq (4u)^{-2}$. Hence if we let $X_{j+1} = \sum_{t=1}^{nu} X_t^{j+1}$, we have

$$\begin{aligned} \Pr(X_{j+1} \geq (2nu)(4u)^{-2}) &\leq \Pr(B(nu, (4u)^{-2}) \geq (2nu)(4u)^{-2}) \\ &\leq e^{-(nu)(4u)^{-2}/3}. \end{aligned}$$

The last part of this equation follows from a standard Chernoff bound (e.g., [13][Theorem 4.4]).

In turn, $S_{j+1}(t)$ can only increase when $P(t) \geq j+1$, and hence, with high probability $S_{j+1}(nu) \leq (2nu)(4u)^{-2}$, so $s_{j+1}(nu) \leq (4u)^{-1}/2$. (This bound is weaker than just taking $s_{j+1}(nu) \leq s_j(nu)(4u)^{-2}$, but will suffice for our induction.)

Let X_{j+k} similarly be the number of steps where $P(t) \geq j+k$. For $k \geq 1$, let $\gamma_k = (4u)^{-2^{k-1}}$; let $\gamma_0 = (4u)^{-2}$. We inductively show that with high probability $s_{j+k}(nu) \leq \gamma_k$ by showing $X_{j+k} \leq n\gamma_k$, as X_{j+k} clearly bounds the number of keys with age $j+k$ placed in the table. To take care of the conditioning issues arising in the induction, we establish a simpler but useful dominating process. Consider a process C with states $j, \dots, j+k$ that initially starts at state j and transitions from state a to b with probability $q(a, b)$ at each time step as follows:

$$q(a, b) = \begin{cases} s_a(t) & \text{for } b = a + 1; \\ s_{j+k}(t) & \text{for } a = b = j + k; \\ s_b(t) - s_{b+1}(t) & \text{for } j < b \leq a \leq j + k \\ & \text{(except as above);} \\ 1 - s_{j+1}(t) & \text{for } b = j. \end{cases}$$

That is, our process C behaves like the age of the key being placed in our original chain, except that the state simply returns to state j as a default (instead of some smaller age) and the maximum state we are concerned with is $j+k$. Because of this, the number of steps where C reaches state $j+k$ over nu times steps stochastically

dominates the number of times the original process has $P(t) \geq j+k$. Also, in order to have for $P(t+k) \geq j+k$ in our original process, the key being placed at time t must have reached age j , and then, over the next k steps, the ages of the cells chosen must be at least $j, j+1, \dots, j+k-1$. Similarly, for process C to reach state $j+k$ the states over the last k steps must successively increment by 1 through $j, j+1, \dots, j+k-1$.

Let Y_{j+k} be the number of steps where C is in state $j+k$ when starting from state j over nu steps. By the dominance of chain C ,

$$\Pr(X_{j+k} \geq \gamma_k) \leq \Pr(Y_{j+k} \geq \gamma_k).$$

In bounding the quantity on the right hand side, we will want to focus on the high probability the case that the $s_{j+i}(t)$ values are well-behaved. In particular, let \mathcal{E}_k be the event that $s_{j+i}(t) \leq \gamma_i$ for $1 \leq i \leq k$ and all $t \leq nu$. Note above we have shown that \mathcal{E}_1 holds with probability $1/n^2$; as part of our proof we show inductively that \mathcal{E}_i holds with probability i/n^2 for i up to a value k^* where $\gamma_{k^*} \leq n^{-2/5}$. Our argument will show that $k^* = \log \log n + O(1)$. At that point we switch to an explicit argument to finish the analysis.

To bound Y_{j+k} , we must consider yet another process \hat{C} , which dominates C conditioned on \mathcal{E}_{k-1} . Specifically, for process \hat{C} , again the states are $j, \dots, j+k$, but the transition from a to b occurs with probability $\hat{q}(a, b)$ at each time step as follows:

$$\hat{q}(a, b) = \begin{cases} \gamma_{a-j} & \text{for } b = a+1; \\ \gamma_k & \text{for } a = b = j+k; \\ \gamma_{b-j} & \text{for } j < b \leq a \leq j+k \\ & \text{(except as above);} \\ \text{all remaining} & \\ \text{probability} & \text{for } b = j. \end{cases}$$

Note $\hat{q}(a, b)$ defines a proper probability distribution. Let Z_{j+k} be the number of steps where \hat{C} is in state $j+k$ when starting from state j over nu steps. Then, conditioned on \mathcal{E}_{k-1} , Z_{j+k} stochastically dominates Y_{j+k} by a simple coupling, since $\hat{q}(a, b) \geq q(a, b)$ whenever $b > j$.

The benefit working with \hat{C} is that the behavior is fixed over time, while C depends on the $s_j(t)$ values. This helps us as follows. We now bound Z_{j+k} using a martingale approach. Let A_i be a binary random variable that is 1 if and only if the process \hat{C} is in state $j+k$ at the i th step, and 0 otherwise. We can use a standard Doob martingale (see e.g. [13][Chapter 12.1]), where $Z_{j+k} = \sum_{i=1}^{nu} A_i$. Note that the A_i are not independent; however, the outcome of each time step of the process affects at most $2k$ values of A_i , and hence the outcome at each time step changes A by at most $2k$.

Hence, using the standard Azuma-Hoeffding inequality (see e.g. [13][Theorem 12.4]), we have

$$\begin{aligned} \Pr(Z_{j+k} \geq 2E[Z_{j+k}]) &\leq \Pr(|Z_{j+k} - E[Z_{j+k}]| \geq E[Z_{j+k}]) \\ &\leq 2\exp(-E[Z_{j+k}]^2/(8nk^2)). \end{aligned}$$

Again, to reach state $j+k$ the states over the last k steps must successively increment by 1 through $j, j+1, \dots, j+k-1$. Hence

$$E[Z_{j+k}] \leq (nu) \prod_{\ell=0}^{k-1} \gamma_{j+\ell} = \frac{n}{4} (4u)^{-2^{k-1}}.$$

Now as stated before let k^* be the smallest value so that $\gamma_{k^*} \leq n^{-2/5}$. Note from the definition of γ_k that $k^* = \log \log n + O(1)$. It follows that, as long as $k < k^* < \log n$, so that $E[Z_{j+k}] \geq n^{3/5}$,

$$\Pr(Z_{j+k} \geq n\gamma_k) \leq \Pr(Z_{j+k} \geq 2E[Z_{j+k}]) \leq 2e^{-n^{1/5}/(8 \log^2 n)} \ll 1/n^2$$

for sufficiently large n . In particular, note that, inductively,

$$\begin{aligned} \Pr(Y_{j+k} \geq \gamma_k) &\leq \Pr(\neg \mathcal{E}_{k-1}) + \Pr(Y_{j+k} \geq \gamma_k \mid \mathcal{E}_{k-1}) \\ &\leq \frac{k-1}{n^2} + \Pr(Z_{j+k} \geq \gamma_k) \\ &\leq \frac{k-1}{n^2} + \frac{1}{n^2} = \frac{k}{n^2}. \end{aligned}$$

Also

$$\begin{aligned} \Pr(\neg \mathcal{E}_k) &\leq \Pr(\neg \mathcal{E}_{k-1}) + \Pr(\neg \mathcal{E}_k \mid \mathcal{E}_{k-1}) \\ &\leq \frac{k-1}{n^2} + \frac{1}{n^2} = \frac{k}{n^2}. \end{aligned}$$

Now for k^* , we follow the same argument to show that

$$\begin{aligned} \Pr(Z_{j+k^*} \geq 2n^{3/5}) &\leq \Pr(|Z_{j+k^*} - E[Z_{j+k^*}]| \geq n^{3/5}) \\ &\leq 2e^{-n^{1/5}/(8 \log^2 n)} < 1/n^2 \end{aligned}$$

for sufficiently large n . Hence, if we let \mathcal{F} be the event that $s_{j+k^*}(t) \leq 2n^{-2/5}$, we find again

$$\begin{aligned} \Pr(\neg \mathcal{F}) &\leq \Pr(\neg \mathcal{E}_{k^*-1}) + \Pr(\neg \mathcal{F} \mid \mathcal{E}_{k^*-1}) \\ &\leq \frac{k^*-1}{n^2} + \frac{1}{n^2} = \frac{k^*}{n^2}. \end{aligned}$$

Now, finally, we note that for any key to have age $j+k^*+3$, the Robin Hood hashing scheme must choose 3 cells with load at least $j+k^*$ over sequential steps. Conditioned on \mathcal{F} , the probability that this event, which we denote by \mathcal{H} , occurs is at most $(2n^{-2/5})^3 \cdot (un) = o(1)$. We conclude

$$\begin{aligned} \Pr(M_n \geq j+k^*+3) &\leq \Pr(M_n \geq j+k^*+3 \text{ and } \mathcal{G}) \\ &\quad + \Pr(M_n \geq j+k^*+3 \text{ and } \neg \mathcal{G}) \\ &\leq \Pr(\mathcal{H}) + \Pr(\neg \mathcal{G}) \\ &\leq \Pr(\mathcal{H} \mid \mathcal{F}) + \Pr(\neg \mathcal{F}) + \Pr(\neg \mathcal{G}) \\ &\leq (2n^{-2/5})^3 \cdot (un) + \frac{k^*}{n^2} + o(1) = o(1). \end{aligned}$$

Here $j + k^* + 3$ is $\log \log n + C$ for a constant C , so the theorem is proven. \square

3.4 Implications for Unsuccessful Search Times

For insertions-only tables, there is a simple optimization that speeds up unsuccessful searches over standard search for Robin Hood hashing. If the key being searched for is at the i th position in its probe sequence, and the probe yields a cell with a key with age strictly less than i , then the key cannot be in the table. This is because if the key were in the table it would have replaced the younger key in this cell. Making use of this fact allows us to short-circuit an unsuccessful search early, before reaching an empty cell. With this optimization, the probability that an unsuccessful search takes at least j probes (up to the longest probe sequence in the system) is $\prod_{k=1}^{j-1} s_k$, since on the k th probe it would need to find a cell with age at least k to continue. In the asymptotic limit, we have that the expected number of probes for an unsuccessful search is thus

$$\sum_{j=1}^{\infty} \prod_{k=1}^{j-1} s_k = \sum_{j=1}^{\infty} j p_j.$$

That is, the expected number of probes for an unsuccessful search is just the expected age according to the equilibrium distribution given by the p_j at the final load α .

This brings up a point that we have not mentioned previously. For Robin Hood hashing, it is useful to keep track of the age of each key with each key. This may require a small number of extra bits per cell, which is not unreasonable if keys are large. As we will see, even for high loads, 3 bits may be sufficient, and 4 bits handles most insertion-only cases in practice. Of course one can always re-derive the age of a key on the fly by re-computing hash values from the key, but we expect keeping the age of the key would prove more efficient.

4 Handling Deletions

In this section, we demonstrate that we can analyze a standard model for deletions whereby we load a number of keys into the system, and then alternate between deleting a key chosen uniformly at random from the table, and inserting new fresh key. Our goal is to use the differential equations analysis to consider the long-term behavior and the steady state (if one exists) of such a system. We show that we can analyze the deletion method described in Celis's thesis [2], which is the only deletion scheme we know of previously proposed for Robin Hood hashing. We then introduce and analyze another deletion scheme that we believe could be more suitable in practice. Further discussion

comparing the two approaches with simulations appears in the appendix. Our scheme is much simpler, and interestingly, the equations we derive to model their behavior are simpler as well.

4.1 Deletions with Tombstones

Celis suggests a scheme for deletions based on a standard approach of using tombstone entries. Cells with deleted keys are marked; such marked entries are called tombstones. For insertion and search purposes deleted entries are treated the same as non-deleted entries, except that when a deleted key is replaced by younger key on an insertion, the deleted key can be discarded and does not need to be put back in the table. As an optimization, if the age of a key being inserted is equal to the age of a key in a deleted entry, it can replace the deleted entry. This maintains the property that if a key is placed in the cell given by the i th entry in its probe sequence, then for $1 \leq j < i$ there is a key (which may be a deleted key in a tombstone cell) with age at most j in the cell corresponding to the j th entry of the key's probe sequence.

This scheme has the obvious problem that the ages of keys in the table can only become larger. Hence, there is no "steady state" for the age of keys. This causes multiple difficulties in practical implementations, which must monitor the minimum and maximum ages in the system and carefully adjust the algorithm accordingly to avoid large search times. We therefore emphasize that we do not believe this deletion scheme is desirable. However, before suggesting an alternative scheme, we show here that we can still model its behavior.

We model this deletion strategy in the setting where we first load the system to a load α , and then alternate between deleting a random key and inserting a new key. Entirely similar analysis can be used for the setting where elements in the system all have an exponentially distributed lifetime with a fixed mean and arrivals of new keys form a Poisson process of a given rate, in which case α represents the equilibrium load of the system. This is a natural model for studying deletions; for example, it was used in Celis's thesis, as well as other works. While this model does not capture possible worst-case behaviors, and in particular does not handle re-insertion of deleted keys, it remains a useful model for understanding basic performance. Of course, it also fits well with our analysis: we start by running the original set of differential equations to load α , and then from the resulting state we run a new set of differential equations that take deletions into account.

We first describe changes to how we think about the state space. Our level process will require an additional state, call it state 0, which corresponds to the state when

we perform a deletion. Also, while we use s_i to again represent the fraction of cells in the table that have an *undeleted* key of age at least i , we also need to track the fraction of cells in the table that are tombstones containing *deleted* keys of age at least i . Let us refer to these as u_i .

We first consider the level process. Let us think in unscaled time steps. Let $q(t)$ be the probability of being in state 0 – that is, that we are about to perform a deletion – and as before for $i \geq 1$ let $p_i(t)$ be the probability that we are trying to place a key of age at least i in the table. The state of the level process is a Markov chain assuming that the state of the table is fixed at certain values s_i and u_i .

The equations for $p_i(t)$ and $q(t)$ are as follows:

$$\begin{aligned} p_1(t+1) &= 1 - \sum_{j=1}^{\infty} (p_j(t) - p_{j+1}(t))(1 - s_1 - u_{j+1}); \\ p_i(t+1) &= p_{i-1}(t)s_{i-1} + \sum_{j=i-1}^{\infty} (p_j(t) - p_{j+1}(t))u_{j+1}; \\ q(t) &= \sum_{j=1}^{\infty} (p_j(t) - p_{j+1}(t))(1 - s_1 - u_{j+1}). \end{aligned}$$

(The second line holds for $i \geq 2$.) Our equation for p_i now has an additional term that takes into account that a key will be placed in a cell with a deleted key when the deleted key's age is less than or equal to the age of key being placed. Similarly, the equation for q_j is based on summing over each possible age j of the key being placed the probability that it is placed successfully. This gives the equilibrium equations:

$$\begin{aligned} p_1 &= 1 - \sum_{j=1}^{\infty} (p_j - p_{j+1})(1 - s_1 - u_{j+1}); \\ p_i &= p_{i-1}s_{i-1} + \sum_{j=i-1}^{\infty} (p_j - p_{j+1})u_{j+1}, \quad i \geq 2; \\ q &= \sum_{j=1}^{\infty} (p_j - p_{j+1})(1 - s_1 - u_{j+1}). \end{aligned}$$

We now turn to equations for s_i and u_i . We find

$$\begin{aligned} \frac{ds_i}{dt} &= \sum_{j=i}^{\infty} (p_j - p_{j+1})(1 - s_i - u_{j+1}) - \frac{qs_i}{s_1}; \\ \frac{du_i}{dt} &= \frac{qs_i}{s_1} - \sum_{j=i}^{\infty} p_j(u_j - u_{j+1}). \end{aligned}$$

Note that, in the above, the values s_1 in the denominator of the terms of the form $\frac{qs_i}{s_1}$ can be replaced by α , since s_1 is α once we start performing deletions under our model.

Unfortunately, this system appears inherently infinite-dimensional. A tactical approach is to modify the system to consider all keys with age greater than or equal to some large constant L to be treated the same; they all are considered to be of age L . In such cases, keys being placed with age at least L can only replace other keys with age at least L if they have been deleted; in all other cases, the process behaves the same. Intuitively, over any finite time (corresponding to cn balls), one can choose L large enough so that the gap between the systems is small enough that it can be ignored. Theorem 2.1 can then be applied to the finite-dimensional system.

Alternatively, for another finite dimensional system where Theorem 2.1 would apply, one could truncate the above equations to correspond to a system where any key with age larger than L simply thrown out of the system (or, equivalently, put in a separate stash). When such a key occurs, we to keep the load constant at α , we assume the misbehaving key is removed and replaced by a fresh key with age 1. Again, by choosing a sufficiently high load threshold, the effect of the truncation should be arbitrarily small. That is, the fraction of discarded keys can be made less than γ for any constant $\gamma > 0$; in practice, for finite systems, for a suitably large truncation threshold with high probability the number of such stashed keys is 0. Also, again Theorem 2.1 can be applied to the finite-dimensional system.

The two families of equations for these limiting systems are both governed by the following equations:

$$\begin{aligned} p_1 &= 1 - \sum_{j=1}^L (p_j - p_{j+1})(1 - s_1 - u_{j+1}); \\ q &= \sum_{j=1}^L (p_j - p_{j+1})(1 - s_1 - u_{j+1}); \\ \frac{ds_i}{dt} &= \sum_{j=i}^L (p_j - p_{j+1})(1 - s_i - u_{j+1}) - \frac{qs_i}{s_1}; \\ \frac{du_i}{dt} &= \frac{qs_i}{s_1} - \sum_{j=i}^L p_j(u_j - u_{j+1}). \end{aligned}$$

the difference between the two systems lies solely in the equation for p_i . For the case where all keys of age at least L are treated equally, we have

$$p_i = p_{i-1}s_{i-1} + \sum_{j=i-1}^{L-1} (p_j - p_{j+1})u_{j+1}.$$

For the case where items of age greater than L are removed and a new key is inserted in its place, we have

$$p_i = p_{i-1}s_{i-1} + \sum_{j=i-1}^{L-1} (p_j - p_{j+1})u_{j+1} - p_L(s_L + u_1 - u_L).$$

Note that for convenience in the above we assume $p_{L+1} = u_{L+1} = s_{L+1} = 0$; fixing these values at zero simplifies the description of the equations.

Ideally, one would formalize the behavior of the infinite-dimensional system by bounding it between two finite-dimensional systems, using stochastic majorization. This is the approach used in [11]. However, to this point we have not found a suitable majorization argument for bounding this variation of the Robin Hood strategy. We leave it as an open problem to formally show that the infinite system accurately models Robin Hood hashing under this deletion strategy for any finite length of time. This conjectured behavior appears reliable based on simulations, where the differential equations are extremely accurate, and we have shown finite-dimensional variations where Theorem 2.1 applies directly. Hence for the practical purposes of generating accurate numerical predictions of performance, we are well placed.

4.2 Deletions without Tombstones

We now suggest and analyze a simpler scheme that we believe remains effective based on our analysis and simulations. Deleted entries are simply deleted from the table; no tombstones are used. A problem with this approach is that one can no longer use an empty cell as a stopping criterion for an unsuccessful search. Similarly, we no longer have the property that on a search an occupied cell must have a key of age at most i on the i th probe, or we can declare the search unsuccessful. The only way to cope with unsuccessful searches is to keep track of largest age of any key currently in the system. This can be done by having the table keep counters of the number of keys of each age.

As before, let us first consider the level process. As a reminder we work in the model where we load the system to a load α , and then alternate between deleting a random key and inserting a new key. Our level process will therefore require an additional state, call it state 0, which corresponds to a deletion. Let us think in unscaled time steps. Let $q(t)$ be the probability of being in state 0 – that is, that we are about to perform a deletion – and as before for $i \geq 1$ let $p_i(t)$ be the probability that we are trying to place a key of age at least i in the table. The state is again a Markov chain, assuming that the state of the table is fixed at certain values s_i .

First, let us consider $q(t)$. When we are placing an item, we complete the placement with probability $1 - s_1$, the probability of finding a cell without a key (either from deletion or from being empty). Note that

$1 - s_1$ is $1 - \alpha$ based on our model. Hence

$$q(t+1) = p_1(t)(1 - \alpha).$$

However, $p_1(t) = 1 - q(t)$. Substituting gives

$$q(t+1) = (1 - q(t))(1 - \alpha).$$

Letting q be the equilibrium probability for the chain for $q(t)$, we have

$$q = \frac{1 - \alpha}{2 - \alpha}.$$

Note that this gives the equilibrium probability

$$p_1 = \frac{1}{2 - \alpha}.$$

Finally, as with the original Robin Hood process, we have for $i \geq 2$ that

$$(4.8) \quad p_i = p_{i-1}s_{i-1}.$$

With this we turn our attention to the limiting equations for the s_i . Note that s_1 increases whenever an available cell is found for a placement, and s_1 decreases whenever a deletion occurs. Hence

$$(4.9) \quad \frac{ds_1}{dt} = p_1(1 - s_1) - q.$$

Note that we have $\frac{ds_1}{dt} = 0$ when $s_1 = \alpha$, so this equation is consistent with our model.

For s_i when $i > 1$, Equation (4.9) generalizes to

$$\frac{ds_i}{dt} = p_i(1 - s_i) - q(s_i/s_1).$$

That is, a cell containing a key with age at least i is created whenever the age of the key being placed is at least i and the probe sequence finds either an empty cell or a cell containing a key with age less than i . A cell containing a key with age at least i is removed whenever a deletion occurs, and that deletion is for a cell holding a key of age at least i , which occurs with probability s_i/s_1 . Assuming we start with alternating deletions and insertions when $s_1 = \alpha$, we can write for $i \geq 1$:

$$(4.10) \quad \frac{ds_i}{dt} = p_i(1 - s_i) - q(s_i/\alpha).$$

Note that, under this model, we again have that the s_i term depends only on values of p_j and s_j with $j \leq i$; hence we can apply Theorem 2.1 over finite time intervals to the truncated family of equations up to $i \leq L$ for any constant L to obtain accurate values for the limiting system. This gives the following result.

THEOREM 4.1. *For any fixed constant L and any $\alpha < 1$, for $i \leq L$, let $s_i(t)$ be the solution for the s_i at time $t \geq \ln \frac{1}{1-\alpha}$ from the family of differential equations given by Equation (4.10) above, when starting at time $t_0 = \ln \frac{1}{1-\alpha}$ with $s_i(t_0)$ obtained from the differential equations given by Equation (3.4). For $1 \leq i \leq L$, let $X_{i,n,t}$ be the random variable denoting the fraction of cells with keys of age at least i using Robin Hood hashing without tombstones after loading the initially empty table to load α , and then alternating insertions of new items and deletions of items independently and uniformly at random, until $(t-t_0)n$ steps have occurred, where a step is either a deletion or an attempted insertion. Then for any $\epsilon > 0$, for sufficiently large n ,*

$$\Pr(|X_{i,n,t} - s_i(t)| > \epsilon) \leq C_1 \exp(-nC_2(\epsilon)),$$

where C_1 is a constant that depends on L and α , and $C_2(\epsilon)$ is a constant that depends on L , α , and ϵ .

Here we have left the theorem statement in terms of time, but as with Theorem 3.1, this is not important. We could have written the statement in terms of the number of pairs of insertions and deletions, as each insertion takes on average $1/\alpha$ steps and each deletion takes 1 step, so time $t = t_0 + z$ in the fluid limit corresponds to $zn\alpha/(1+\alpha)$ insertion-deletion pairs.

For this model, we find that there is a unique equilibrium distribution for the underlying Equations (4.10) and (4.8); this gives us an idea as to the long term performance of this approach. We denote the equilibrium values by s_i^* and p_i^* . In equilibrium, using $s_1 = \alpha$ always, we find that $ds_i/dt = 0$ gives the following reasonable equation:

$$(4.11) \quad s_i^* = \frac{p_i^*}{p_i^* + \frac{1-\alpha}{\alpha(2-\alpha)}}.$$

Equation (4.11) along with Equation (4.8) can be used to show that the s_i^* again decrease double exponentially at the equilibrium given by the family of Equations (4.10).

THEOREM 4.2. *In the asymptotic model for an infinite Robin Hood hash table with load factor $\alpha < 1$ and alternating deletions, for sufficiently large constants i , the value of s_i^* at the equilibrium point where $ds_i/dt = 0$ everywhere satisfies*

$$s_i^* \leq c_1 c_2^{2^{i-c_3}}$$

for some constants $c_1, c_3 > 0$ and $c_2 < 1$ that may depend on α .

Proof. In what follows let s_i^* and p_i^* refer to their values in equilibrium. Let $z = \frac{1-\alpha}{\alpha(2-\alpha)}$. From Equation (4.11),

$s_i^* < p_i^*/z$. If $z \geq 1$, using this and $p_i^* = s_{i-1}^* p_{i-1}^*$, we can induct to find

$$s_i^* \leq \left(\frac{\alpha}{z(2-\alpha)} \right)^{2^{i-2}}.$$

For case $z < 1$, we note the tails of the p_i^* must decrease geometrically, since a key is placed in an empty cell or a cell with a deleted item with probability at least $1-\alpha$ at each step. Let j be the smallest value such that $p_j^* \leq z^2$. Then inductively we find

$$s_{j+k}^* \leq z^{2^k}.$$

□

Finally, we note that the differential equations are easily shown to converge to the equilibrium point by induction on i for the s_i . Equation (4.10) shows that

$$\frac{ds_i}{dt} = p_i - s_i \left(p_i + \frac{1-\alpha}{\alpha(2-\alpha)} \right).$$

We have $s_1 = \alpha$ and $p_1 = 1/(2-\alpha)$; these values are fixed at their equilibrium values. Inductively, we find that p_i , which depends on values s_j for $j < i$, converges to its equilibrium. Suppose $p_i \in (p_i^* - \epsilon, p_i^* + \epsilon)$ for some $\epsilon > 0$. Then s_i is decreasing if

$$s_i \geq \frac{p_i^* + \epsilon}{p_i^* + \epsilon + \frac{1-\alpha}{\alpha(2-\alpha)}} = s_i^* + \epsilon_1,$$

and similarly s_i is increasing if $s_i \leq s_i^* - \epsilon_2$ for ϵ_1, ϵ_2 that go to 0 as ϵ goes to 0. It follows that after long enough periods of time the differential equations are close to the equilibrium point. In simulations, we see the convergence happens suitably fast to accurately describe performance.

5 Conclusion

We have shown how to use the framework of Markov chains, specifically via what is commonly called fluid limit analysis or the mean-field approach, to analyze Robin Hood hashing. In particular, we have shown that for Robin Hood hashing the analysis naturally requires the use of an additional level process. Besides providing a new way of gaining insight into previous results, we have shown that our methods lead to a simple recurrence describing the equilibrium behavior of Robin Hood hashing under a natural deletion model when not using tombstones.

Robin Hood hashing appears to perform essentially the same whether using probe sequences based on double hashing and random hashing. Proving this seems a worthwhile open question. Relatedly, the recent work

of [9] applies fluid limit analysis to show that double hashing yields the same behavior as fully random hashing under the “balanced allocation” paradigm. Alternatively, one could try to extend the approach of [6] used for standard open addressing hashing, as done in [10] for the “balanced allocation” setting.

References

- [1] Y. Azar, A. Broder, A. Karlin, E. Upfal. Balanced allocations. *SIAM Journal on Computing*, 29(1):180-200, 1999.
- [2] P. Celis. Robin Hood Hashing. Ph.D. thesis, University of Waterloo, 1986. Available as Technical Report CS-86-14, Computer Science Department, University of Waterloo.
- [3] P. Celis, P. Larson, and I. Munro. Robin hood hashing. In *Proceedings of the 26th Annual Symposium on Foundations of Computer Science*, pp. 281-288, 1985.
- [4] L. Devroye, P. Morin, and A. Viola. On worst-case robin hood hashing. *SIAM Journal on Computing*, 33(4):923-936, 2004.
- [5] S. Janson. Individual displacements for linear probing hashing with different insertion policies. *ACM Transactions on Algorithms*, 1(2):177-213, 2005.
- [6] G. Lueker and M. Molodowitch. More analysis of double hashing. *Combinatorica*, 13(1):83-96, 1993.
- [7] M. Luczak, and J. Norris. Averaging over fast variables in the fluid limit for Markov chains: application to the supermarket model with memory. *The Annals of Applied Probability*, 23(3):957-986, 2013.
- [8] M. Mitzenmacher. The power of two choices in randomized load balancing. Ph.D. thesis, U.C. Berkeley, 1996.
- [9] M. Mitzenmacher. Balanced allocations and double hashing. In *Proceedings of the 26th ACM Symposium on Parallel Algorithms and Architectures*, pp. 331-342, 2014.
- [10] M. Mitzenmacher. More analysis of double hashing for balanced allocations. In *Proceedings of the 13th Workshop on Analytic Algorithms and Combinatorics*, 2016.
- [11] M. Mitzenmacher, B. Prabhakar, and D. Shah. Load balancing with memory. In *Proceedings of the 43rd Annual Symposium on Foundations of Computer Science*, pp. 799-808, 2002.
- [12] M. Mitzenmacher, A. Richa, and R. Sitaraman. *The Power of Two Choices: A Survey of Techniques and Results*, edited by P. Pardalos, S. Rajasekaran, J. Reif, and J. Rolim. Kluwer Academic Publishers, Norwell, MA, 2001, pp. 255-312.
- [13] M. Mitzenmacher and E. Upfal. Probability and computing: Randomized algorithms and probabilistic analysis. Cambridge University Press, 2005.
- [14] A. Shwartz, and A. Weiss. Large deviations for performance analysis: Queues, communication and computing. Chapman & Hall, 1995.
- [15] A. Viola. Analysis of hashing algorithms and a new

mathematical transform. Ph.D. thesis, University of Waterloo, 1995. Available as Technical Report CS-95-50, Computer Science Department, University of Waterloo.

- [16] A. Viola. Exact distribution of individual displacements in linear probing hashing. *ACM Transactions on Algorithms*, 1(2):214-242, 2005.

Appendices

Appendix A: Proof of Theorem 3.1

Proof. The result follows from Theorem 2.1. While Theorem 2.1 is stated in terms of time instead of load, this difference is not consequential, as we explain subsequently. Let $Y_{i,n}$ be the random variable denoting the fraction of cells with keys of age at least i using Robin Hood hashing after $-n \ln(1 - \alpha)$ unscaled time steps in a hash table with n cells. Then Theorem 2.1 gives us that

$$\Pr(|Y_{i,n} - s_i(-\ln(1 - \alpha))| > \epsilon) \leq C_3 \exp(-nC_4(\epsilon)),$$

where C_3 is a constant that depends on K and α , and $C_4(\epsilon)$ is a constant that depends on K , ϵ , and α . Note that this depends on our restriction of the system to be finite dimensional, and the fact that the evolution of the s_i for $i \leq K$ only depends on the values s_1, s_2, \dots, s_K . (Again, this is another reason to write equations in terms of the tails of the loads.) The conditions of Theorem 2.1 are easily checked. In particular, we have noted the Markov process over levels is ergodic for any given s_i values. For the second condition, the transition rates $\nu_i(\bar{x}; m)$ are finite size polynomials of load vector \bar{x} . This implies that $\log \nu_i(\bar{x}; m)$ is Lipschitz continuous in coordinates of \bar{x} . In particular, the bound $x_i \leq 1$ gives an upper bound. If $\nu_i(\bar{x}; m) = 0$, then it means the transition is absent and we neglect it. Otherwise $x_i \geq 1/n_0$, giving a lower bound on the $\log \nu_i(\bar{x}; m)$. This completes the check of condition 2.2.

Now note that $\hat{s}_i(\alpha) = s_i(-\ln(1 - \alpha))$ in the limiting system. We find $X_{i,n}$ and $Y_{i,n}$ differ by $o(1)$ terms with high probability, and in fact

$$\Pr(|Y_{i,n} - X_{i,n}| > \gamma) \leq C_5 \exp(-nC_6(\gamma)),$$

where C_5 is a constant that depends on K and α and $C_6(\gamma)$ is a constant that depends on K , γ , and α . We sketch the reasoning: consider the coupling where we perform the Robin Hood process for the maximum of $-n \ln(1 - \alpha)$ time steps and the number of time steps to reach load α . The load after $-n \ln(1 - \alpha)$ time steps will be $\alpha \pm o(1)$ with high probability, by standard martingale arguments; alternatively, the number of time steps to reach load α is $-n \ln(1 - \alpha) + o(n)$ with high probability by standard Chernoff-type bounds, since

the number of time steps to place each key is an independent geometric random variable with bounded mean. The theorem statement holds by summing the probabilities $\Pr(|Y_{i,n} - s_i(-\ln(1-\alpha))| > \epsilon/2)$ and $\Pr(|Y_{i,n} - X_{i,n}| > \epsilon/2)$. \square

Appendix B: Simulations for Insertions Only

In this section we provide simulation results. These results serve the dual purpose of demonstrating the effectiveness of Robin Hood hashing and verifying our analysis. We note that simulation results were also presented in [2], and one might look there for further discussion on the effectiveness of Robin Hood hashing. Here, the simulation results are presented for completeness, to provide a high-level verification of the utility of the theoretical framework.

Table 1 show results with a load α of 0.95 on the hash table. The fraction of keys in the table with a given age (up to 7) are given. The results from the differential equations were calculated using the standard Euler's method with discrete time steps of length 10^{-6} ; that is, we calculate successive estimates of the variables from the differential equations using the derivative at the current values and advance time in steps of 10^{-6} . We also calculate the results from the recurrence of Theorem 3.2. For our simulation results, the probe sequence for all of the elements were determined using the pseudo-random generator drand48; all trials were performed in a single seeded run. For the simulations, the expression $x \pm y$ refers to the average x and standard deviation y over 1000 trials. Unsurprisingly, the differential equations and the results from Theorem 3.2 agree quite closely, with the discrepancy explained by our calculation method for the differential equations. The theoretical results match the simulations very closely.

These results also show the potential effectiveness of Robin Hood hashing. Even at a load of 0.95, the maximum age over these simulations was 7. The average number of probes for a successful search (of a random key in the table) is approximately 3.15; the average number of probes for an unsuccessful search is approximately 3.59. As pointed out by Celis [2], there are further possible ways to speed up searches by using procedures other than the standard search procedure. For example, since most keys have age 3 or 4 under this load, one can start the search with the third and fourth entries of the probe sequence to reduce the expected number of cells examined for a successful search. Note this would be possible given a hash function in the form $h: K \times \mathbb{N} \rightarrow [0, n-1]$, so we can examine the i th entry in the probe sequence directly.

It is worth noting (as was noted in [2] as well)

that the results appear essentially unchanged even if double hashing is used instead of (our proxy for) fully random hashing. (In double hashing, we choose a starting point a and an offset b that is relatively prime to the table size and our probe sequence is given by $a, a+b, a+2b, \dots$, where the values in the probe sequence are taken modulo the size of the hash table.) Table 2 shows a representative example. We are not sure yet how to prove this, although we suspect that theoretical techniques that have shown double hashing has the same performance as fully random hashing in other settings may apply (e.g., [6, 9]). The challenge lies in accounting for the ages of placed keys in such an analysis.

Appendix C: Simulations with Deletions

We first consider the setting with tombstones. Here our goal in the simulation is simply to show that the proposed differential equations accurately model the actual system for finite periods of time correctly. In order to keep tables reasonably sized, the result of Table 3 shows results with a load α of 0.9; here we load the table with $0.9n$ items, and then alternately delete and insert items until $2n$ items have been placed. For the differential equations we again use Euler's method with discrete time steps of length 10^{-6} . Here we show the fraction of keys in the system (not including tombstones) by age. The results are very accurate, and begins to show the effect of using tombstones; the keys of ages 1 and 2 are vanishing, and the ages of the keys in the cells are increasing over time. If we continued the simulation further, we would see ages continue to grow well beyond 18.

Table 4 shows results with the same setup for the load and deletion pattern. Here the maximum age does not increase the same way, as noted in our analysis, and the maximum age remains smaller at 12. Again, the differential equations prove highly accurate. In this case, however, we are further interested in the equilibrium distribution of the age as given by Equation (4.11). As a proxy we run the simulations again, but after loading the table with $0.9n$ items, we alternately delete and insert items until $10n$ items have been placed. Table 5 shows these results, compared to the calculated equilibrium distribution from Equation (4.11) (which was derived from the corresponding differential equations). Again, we see that the results match well, showing the utility of the differential equation. Also, as shown Theorem 4.2, we see the probability a cell has a certain age falls very quickly (doubly exponentially) at the tail of the distribution. The average time for a successful search naturally converges to 10. The maximum age, and hence the time for an unsuccessful search, is only around 16.

Key Age	Differential Equations	Celis Theorem	Sims $n = 8192$	Sims $n = 65536$	Sims $n = 524288$
1	0.083458328	0.083458403	$0.083621434 \pm 0.004965813$	$0.083429753 \pm 0.001831579$	$0.083428593 \pm 0.000614445$
2	0.188976794	0.188976856	$0.189157158 \pm 0.008679919$	$0.189075507 \pm 0.003140080$	$0.188945002 \pm 0.001092080$
3	0.323793458	0.323793385	$0.323707145 \pm 0.008193784$	$0.323798214 \pm 0.003014369$	$0.323775391 \pm 0.001046901$
4	0.303363752	0.303363594	$0.302934079 \pm 0.008740126$	$0.303259705 \pm 0.003202237$	$0.303395377 \pm 0.001123541$
5	0.095303269	0.095303242	$0.095385891 \pm 0.009795472$	$0.095321415 \pm 0.003651414$	$0.095341351 \pm 0.001226667$
6	0.005092100	0.005092104	$0.005182087 \pm 0.001441913$	$0.005103134 \pm 0.000523627$	$0.005101511 \pm 0.000174886$
7	0.000012417	0.000012417	$0.000012208 \pm 0.000039393$	$0.000012271 \pm 0.000014530$	$0.000012775 \pm 0.000005037$

Table 1: Results for Robin Hood hashing, both theoretical and from simulations. The simulations are for $\alpha = 0.95$ and 1000 trials.

Key Age	Differential Equations	Celis Theorem	Sims $n = 8192$	Sims $n = 65536$	Sims $n = 524288$
1	0.083458328	0.083458403	$0.083847726 \pm 0.004951031$	$0.083431054 \pm 0.001784843$	$0.083421016 \pm 0.000612585$
2	0.188976794	0.188976856	$0.189757389 \pm 0.008688261$	$0.188963523 \pm 0.003113784$	$0.188964048 \pm 0.001082245$
3	0.323793458	0.323793385	$0.324344385 \pm 0.008153945$	$0.323768419 \pm 0.002991134$	$0.323744102 \pm 0.001028609$
4	0.303363752	0.303363594	$0.302252763 \pm 0.008952148$	$0.303332755 \pm 0.003156650$	$0.303420142 \pm 0.001083855$
5	0.095303269	0.095303242	$0.094714212 \pm 0.009595910$	$0.095374789 \pm 0.003633155$	$0.095344666 \pm 0.001237186$
6	0.005092100	0.005092104	$0.005069262 \pm 0.001446734$	$0.005116963 \pm 0.000521156$	$0.005093440 \pm 0.000182278$
7	0.000012417	0.000012417	$0.000014264 \pm 0.000041576$	$0.000012496 \pm 0.000014498$	$0.000012587 \pm 0.000005136$

Table 2: Results for Robin Hood hashing, both theoretical and from simulations, but here the simulations use double hashing.

Key Age	Differential Equations	Sims $n = 8192$	Sims $n = 65536$	Sims $n = 524288$
1	0.0000000012	0 ± 0	0 ± 0	0 ± 0
2	0.0000000088	0 ± 0	0 ± 0	$0.0000000042 \pm 0.0000000947$
3	0.0000000621	$0.0000004069 \pm 0.0000074176$	$0.0000000509 \pm 0.0000009272$	$0.0000000721 \pm 0.0000003956$
4	0.0000004128	$0.0000001356 \pm 0.0000042869$	$0.0000003730 \pm 0.0000025999$	$0.0000004472 \pm 0.0000009996$
5	0.0000025165	$0.0000016276 \pm 0.0000147681$	$0.0000024584 \pm 0.0000062977$	$0.0000026046 \pm 0.0000023888$
6	0.0000140033	$0.0000168181 \pm 0.0000478803$	$0.0000138008 \pm 0.0000151588$	$0.0000140084 \pm 0.0000054692$
7	0.0000711550	$0.0000779872 \pm 0.0001004372$	$0.0000708691 \pm 0.0000354163$	$0.0000718202 \pm 0.0000124441$
8	0.0003302926	$0.0003438220 \pm 0.0002233071$	$0.0003195890 \pm 0.0000746674$	$0.0003317559 \pm 0.0000275909$
9	0.0014006589	$0.0014721280 \pm 0.0005133678$	$0.0013790479 \pm 0.0001734967$	$0.0014051570 \pm 0.0000624493$
10	0.0054203409	$0.0056228130 \pm 0.0012861569$	$0.0053175036 \pm 0.0004155336$	$0.0054218866 \pm 0.0001541873$
11	0.0190426783	$0.0195947376 \pm 0.0033890440$	$0.0187390051 \pm 0.0011403881$	$0.0190677512 \pm 0.0004144598$
12	0.0596408085	$0.0610249559 \pm 0.0089433574$	$0.0587034689 \pm 0.0029890910$	$0.0597147622 \pm 0.0010595725$
13	0.1576758321	$0.1599762648 \pm 0.0177522046$	$0.1557674884 \pm 0.0061968220$	$0.1578802884 \pm 0.0021976097$
14	0.3056376472	$0.3063086939 \pm 0.0174270189$	$0.3036759011 \pm 0.0063584378$	$0.3057603924 \pm 0.0022285217$
15	0.3239990269	$0.3198004883 \pm 0.0176422773$	$0.3255669696 \pm 0.0058555114$	$0.3238154957 \pm 0.0021086756$
16	0.1187678782	$0.1173343280 \pm 0.0260781256$	$0.1218899664 \pm 0.0094442642$	$0.1185390890 \pm 0.0033043206$
17	0.0079676382	$0.0083817985 \pm 0.0044648123$	$0.0085186328 \pm 0.0015278714$	$0.0079452633 \pm 0.0005104320$
18	0.0000292668	$0.0000429947 \pm 0.0000969077$	$0.0000348751 \pm 0.0000273052$	$0.0000292015 \pm 0.0000087025$

Table 3: Results for Robin Hood hashing with deletions and tombstones, both theoretical and from simulations. Here $2n$ total items are inserted. The simulations are for $\alpha = 0.90$ and 1000 trials.

Key Age	Differential Equations	Sims $n = 8192$	Sims $n = 65536$	Sims $n = 524288$
1	0.0109912456	0.0110180068 \pm 0.0013529380	0.0110012567 \pm 0.0005847770	0.0109806228 \pm 0.0003822734
2	0.0132627846	0.0132962114 \pm 0.0014180089	0.0132489158 \pm 0.0006426987	0.0132499410 \pm 0.0004533081
3	0.0164099523	0.0164749126 \pm 0.0017334129	0.0164142764 \pm 0.0007593122	0.0163919906 \pm 0.0005562663
4	0.0215495612	0.0216120513 \pm 0.0020547157	0.0215255058 \pm 0.0009606507	0.0215314282 \pm 0.0007204324
5	0.0330311741	0.0333027883 \pm 0.0029126093	0.0329123421 \pm 0.0014204389	0.0330133628 \pm 0.0010944805
6	0.0655968369	0.0660263377 \pm 0.0054606729	0.0653358713 \pm 0.0026905826	0.0655498797 \pm 0.0021627231
7	0.1508513719	0.1515972160 \pm 0.0111574845	0.1502017361 \pm 0.0059413695	0.1507750735 \pm 0.0049173364
8	0.2865694087	0.2866796841 \pm 0.0139163312	0.2856982460 \pm 0.0097570156	0.2864050409 \pm 0.0091539804
9	0.2955178737	0.2936264874 \pm 0.0145071670	0.2955954899 \pm 0.0101189923	0.2951344134 \pm 0.0094283266
10	0.1003906169	0.0994804734 \pm 0.0148415190	0.1011552450 \pm 0.0060279724	0.1001607244 \pm 0.0036143267
11	0.0058132004	0.0058682671 \pm 0.0020576713	0.0058955320 \pm 0.0007574272	0.0057927890 \pm 0.0003043712
12	0.0000159736	0.0000185627 \pm 0.0000504047	0.0000165817 \pm 0.0000181370	0.0000157326 \pm 0.0000057864

Table 4: Results for Robin Hood hashing with deletions and no tombstones, both theoretical and from simulations. Here $2n$ total items are inserted. The simulations are for $\alpha = 0.90$ and 1000 trials.

Key Age	Calculated Equilibrium	Sims $n = 8192$	Sims $n = 65536$	Sims $n = 524288$
1	0.0109890110	0.0110295238 \pm 0.0013242152	0.0109825070 \pm 0.0005826258	0.0109803031 \pm 0.0003835962
2	0.0132380001	0.0132215540 \pm 0.0014401913	0.0132458671 \pm 0.0006504940	0.0132260213 \pm 0.0004546164
3	0.0162108987	0.0162228928 \pm 0.0016968246	0.0161892632 \pm 0.0007604040	0.0161916268 \pm 0.0005516735
4	0.0202345136	0.0202062958 \pm 0.0019607504	0.0202353850 \pm 0.0009027583	0.0202174129 \pm 0.0006767551
5	0.0258283516	0.0259177958 \pm 0.0022695870	0.0258402714 \pm 0.0010973218	0.0258078395 \pm 0.0008603925
6	0.0338433436	0.0339240307 \pm 0.0027397431	0.0338269272 \pm 0.0013708900	0.0338119259 \pm 0.0011128289
7	0.0457090363	0.0458778980 \pm 0.0034805126	0.0457426674 \pm 0.0017698065	0.0456839344 \pm 0.0014981448
8	0.0638449846	0.0640809076 \pm 0.0043114573	0.0638339146 \pm 0.0024221701	0.0637959164 \pm 0.0020769152
9	0.0921369579	0.0923216211 \pm 0.0056647485	0.0921366789 \pm 0.0034257988	0.0920613574 \pm 0.0029824930
10	0.1351848968	0.1354930164 \pm 0.0077845001	0.1351204660 \pm 0.0048713224	0.1350763809 \pm 0.0043471712
11	0.1893101510	0.1891220132 \pm 0.0091155314	0.1891978863 \pm 0.0064652559	0.1891614878 \pm 0.0060426133
12	0.2098741222	0.2093249217 \pm 0.0100319761	0.2096149248 \pm 0.0071636039	0.2097165961 \pm 0.0067030989
13	0.1226847741	0.1217393755 \pm 0.0111532706	0.1223768263 \pm 0.0054076900	0.1225484444 \pm 0.0040878863
14	0.0205100133	0.0201279800 \pm 0.0038388175	0.0202694291 \pm 0.0015550304	0.0203358892 \pm 0.0007973781
15	0.0004008004	0.0003907662 \pm 0.0002525555	0.0003878323 \pm 0.0000920074	0.0003857338 \pm 0.0000357398
16	0.0000001447	0.0000004065 \pm 0.0000074139	0.0000001524 \pm 0.0000016004	0.0000001291 \pm 0.0000005412

Table 5: Results for Robin Hood hashing with deletions and no tombstones, both theoretical and from simulations. Here $10n$ total items are inserted, and compared against the calculated equilibrium distribution. The simulations are for $\alpha = 0.90$ and 1000 trials.